


a

SmoIVLA closed-loop control

New observations trigger action-chunk prediction; selected actions are then executed in MuJoCo



b

Inference refinement

$\text{num_steps} = N$

One policy call



More refinements increase computation within each policy invocation

Controls compute per call

c

Prediction horizon

$\text{chunk_size} = H$

One policy output



The full action sequence is generated from the current observation


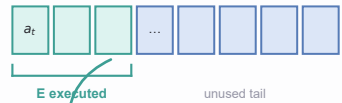
Controls actions predicted per call

d

Execution horizon

$\text{n_action_steps} = E$

Predicted chunk ($E \leq H$)



Controls feedback and replanning frequency